

Table of Contents

1. MCP Server and Client	1
1.1. Module Overview	1
1.2. MCP Server	1
1.2.1. Security	1
1.2.2. Tools	2
1.2.3. Resources	3
1.2.4. Connecting an MCP Client	3
1.3. MCP Client	4
1.3.1. Configuration	4
1.3.2. Using MCP Tools in AI Services	5
1.3.3. Programmatic Access to MCP Clients	5
1.3.4. Dev UI: Testing MCP Clients	6
1.3.5. Optional: Resources as Tools	6
1.3.6. Optional: Headers and Authorization	7
1.4. Agent REST API	7
1.4.1. Endpoints	7
1.4.2. Execute Request	7
1.5. Per-Tenant Agent Configuration	8
1.5.1. Run-As Principal	9
1.6. Agent Configuration API	9
1.6.1. Agent Entity	9
1.6.2. CRUD Endpoints	10
1.6.3. Example: Creating an Agent	10
1.7. Architecture	11
1.8. Ontology Tools	11
1.9. Integrating with Claude	12
1.9.1. Claude Desktop	12
1.9.2. Claude Code (CLI)	12
1.9.3. Claude.ai (Browser) with Custom Connectors	13
1.9.4. Cursor IDE	13
1.10. How MCP Maps to Framework Agent and Tool Concepts	14
1.10.1. Concept Mapping	14
1.10.2. Execution Flow Comparison	14
1.11. End-to-End Example: Claude Querying a Supply Chain Application	15
1.11.1. Prerequisites	15
1.11.2. Step 1: Add the MCP Server Module	15
1.11.3. Step 2: Configure Agent Properties	15
1.11.4. Step 3: Start the Application	16

1.11.5. Step 4: Configure Claude	16
1.11.6. Step 5: Interact with Claude	17
1.11.7. What Happened Under the Hood	18
1.11.8. Adding Ontology Context for Richer Queries.	19
1.12. MCP Server Configuration Reference	19

Chapter 1. MCP Server and Client

The `quantum-mcp-server` module exposes the Query Gateway as a set of [Model Context Protocol \(MCP\)](#) tools and resources, and provides an MCP client for calling external MCP tool providers. This enables AI assistants (Claude Desktop, Cursor, ChatGPT, and others) to discover and invoke your application’s CRUDL operations, browse entity schemas, and receive query-building hints—all through the standard MCP JSON-RPC protocol.

The module also includes a REST-based agent layer (`/api/agent/*`) that mirrors the same tool set for non-MCP integrations.

1.1. Module Overview

`quantum-mcp-server` is a standalone Maven module with three key dependencies:

Dependency	Version	Purpose
<code>quantum-framework</code>	<code>\${quantum.version}</code>	Core framework (QueryGatewayResource, security, Morphia)
<code>quarkus-mcp-server-http</code>	1.9.1	Quarkiverse MCP Server — Streamable HTTP + legacy SSE transport
<code>quarkus-langchain4j-mcp</code>	1.6.0	Quarkiverse MCP Client via LangChain4j — connect to external MCP servers

Add the module to your application’s POM:

```
<dependency>
  <groupId>com.end2endlogic</groupId>
  <artifactId>quantum-mcp-server</artifactId>
  <version>${quantum.version}</version>
</dependency>
```

1.2. MCP Server

The MCP server exposes the Query Gateway as six tools and three resources at the `/mcp` endpoint. MCP clients discover these automatically via the `tools/list` and `resources/list` JSON-RPC methods.

1.2.1. Security

The MCP server and the agent REST API require an authorized user. All endpoints and tools are protected by Quarkus Security.

- **REST API (`/api/agent/*`):** Requires a valid Bearer JWT token in the `Authorization` header.
- **MCP Endpoint (`/mcp`):** The MCP protocol transport (Streamable HTTP or SSE) must include the same authentication mechanism (typically a `Authorization: Bearer <token>` header) as other REST APIs.

You can generate a long-lived service token for use with the MCP server using the [Quantum CLI](#).

1.2.2. Tools

Tools are defined in `McpGatewayTools` using the Quarkiverse `@Tool` and `@ToolArg` annotations. Each tool is protected with `@Authenticated` and delegates to `AgentExecuteHandler`, which routes to the `QueryGatewayResource`—reusing the same security, realm resolution, and query execution as the REST API.



Because the tools route through `QueryGatewayResource`, the data-access tools are row- and field-governed. `query_find`, `query_count`, `query_delete`, `query_deleteMany`, and `query_export` apply the caller's `DataDomain` scope plus the `RuleContext` rule `filterConstraints`, and `query_find` additionally excludes deny-wins `excludedFields` at the MongoDB level. The gateway **fails closed** (HTTP 500 `SecurityContextMissing`) if no security context is established for the request, so an agent never executes an ungoverned query. `query_delete` returns "not found" for an id outside the caller's `DataDomain` rather than revealing that the row exists.

`expand(...)` (aggregation) execution is currently disabled: a `query_find` with `expand(...)` returns HTTP 501 because the aggregation path is not yet governed by row- and field-level security. See [Planner and QueryGateway: row- and field-level governance](#).

Tool Name	Description
<code>query_rootTypes</code>	List available entity types (root types) that can be queried, saved, or deleted
<code>query_plan</code>	Return the query execution plan (FILTER vs AGGREGATION) for a <code>rootType</code> and BIAPI query string
<code>query_find</code>	Execute a BIAPI query and return matching entities (supports pagination, realm override)
<code>query_save</code>	Save (insert or update) an entity by <code>rootType</code>
<code>query_delete</code>	Delete a single entity by its <code>ObjectId</code>
<code>query_deleteMany</code>	Delete multiple entities matching a BIAPI query
<code>query_count</code>	Count entities matching a BIAPI query
<code>query_export</code>	Export entities to a stream (CSV) matching a BIAPI query and column selection
<code>query_import_analyze</code>	Analyze a CSV file for import, returning a preview of rows and identifying potential errors
<code>query_import_rows</code>	List analyzed import rows from a session (supports filtering by error)
<code>query_import_commit</code>	Commit an analyzed import session, persisting the valid rows to the database
<code>query_import_cancel</code>	Cancel an import session and clean up temporary data

Tool parameters

`query_find` accepts these arguments:

Parameter	Type	Description
<code>rootType</code>	string	Entity type simple name or FQCN (use <code>query_rootTypes</code> to discover)
<code>query</code>	string	BIAPI query string (e.g. <code>status:ACTIVE && region:West</code>)
<code>realm</code>	string	Optional tenant realm (defaults to caller's realm)
<code>limit</code>	integer	Optional max results (default 50)
<code>skip</code>	integer	Optional offset for pagination (default 0)

`query_save` accepts `rootType`, `entity` (JSON object matching the schema), and optional `realm`. `query_delete` accepts `rootType`, `id` (ObjectId hex), and optional `realm`. `query_deleteMany` accepts `rootType`, `query`, and optional `realm`.

1.2.3. Resources

Resources are defined in `McpSchemaResources` and `McpQueryHintsResource` using the `@Resource` annotation. MCP clients can read these to populate LLM context with schema information and query-building guidance.

Resource URI	Description
<code>quantum://schema</code>	Lists all available root types with class name, simple name, and collection name
<code>quantum://query-hints</code>	BIAPI query grammar summary, example queries by intent, and tips (expand, wildcards, ontology)
<code>quantum://permission-hints</code>	Permission check/evaluate API summary, area/domain/action mapping, and example check requests

1.2.4. Connecting an MCP Client

Any MCP-compatible client can connect to the `/mcp` endpoint. Example configuration for Claude Desktop (`claude_desktop_config.json`):

```
{
  "mcpServers": {
    "quantum": {
      "url": "http://localhost:8080/mcp"
    }
  }
}
```

For Cursor, add the server URL in Settings > MCP Servers.

Once connected, the client can:

1. Call `tools/list` to discover the six gateway tools
2. Call `resources/list` to discover schema and hint resources
3. Call `resources/read` with `quantum://query-hints` to learn the BIAPI query syntax
4. Call `tools/call` with `query_rootTypes` to see what entity types are available
5. Call `tools/call` with `query_find` to query data

1.3. MCP Client

The MCP client side uses the Quarkiverse LangChain4j MCP extension (`quarkus-langchain4j-mcp`) to connect to external MCP servers and consume their tools. Use the MCP client when you need to:

- Integrate with external tool providers (e.g. Helix MCP, Brain, HelixAI, or any MCP-compliant server)
- Let your AI services discover and invoke tools from those servers at runtime
- Access MCP resources (files, API responses, etc.) programmatically or expose them as tools to the LLM

The client implements the [MCP Specification](#) and works with any compatible MCP server (including the Quantum MCP server on another instance).

1.3.1. Configuration

External MCP connections are configured in `application.properties` using the `quarkus.langchain4j.mcp.<client-name>` prefix. Each named block defines one MCP client (e.g. `helix`, `brain`, `filesystem`).

```
# Streamable HTTP (recommended for remote servers)
quarkus.langchain4j.mcp.helix.transport-type=streamable-http
quarkus.langchain4j.mcp.helix.url=http://helix-mcp.example.com/mcp

# STDIO (for local subprocess servers, e.g. npm-based MCP servers)
quarkus.langchain4j.mcp.brain.transport-type=stdio
quarkus.langchain4j.mcp.brain.command=npx,-y,@brain/mcp-server

# Optional: pass environment variables to the STDIO subprocess
quarkus.langchain4j.mcp.brain.environment.API_KEY=${EXTERNAL_API_KEY}
```

Supported transport types:

Transport	Description
<code>streamable-http</code>	Recommended for remote MCP servers; uses HTTP with the MCP streamable transport.
<code>http</code>	Legacy HTTP/SSE transport. URL typically ends with <code>/mcp/sse</code> .

Transport	Description
<code>stdio</code>	Spawns the MCP server as a subprocess; communication via stdin/stdout. Use <code>command</code> (comma-separated parts) and optionally <code>environment.*</code> for env vars.
<code>websocket</code>	Bidirectional WebSocket. Supported by some servers (e.g. those built with <code>quarkus-mcp-server</code>).

1.3.2. Using MCP Tools in AI Services

Enable tools from one or more MCP servers on an AI service using `@McpToolBox`. The LLM can then discover and call those tools during the conversation.

Single MCP server:

```
import io.quarkiverse.langchain4j.mcp.runtime.McpToolBox;
import dev.langchain4j.service.SystemMessage;
import dev.langchain4j.service.UserMessage;
import io.quarkiverse.langchain4j.RegisterAiService;

@RegisterAiService
public interface MyAiService {

    @SystemMessage("You are a helpful assistant.")
    @McpToolBox("helix")
    String chat(@UserMessage String userMessage);
}
```

Multiple MCP servers: pass an array of client names to use tools from several servers in one service.

```
@McpToolBox({"helix", "brain"})
String chat(@UserMessage String userMessage);
```

All configured MCP servers: omit the value to expose every configured MCP client's tools to the AI service.

```
@McpToolBox
String chat(@UserMessage String userMessage);
```

1.3.3. Programmatic Access to MCP Clients

For fine-grained control (custom tool filtering, non-AI workflows, or resource access), inject an `McpClient` directly with `@McpClientName`:

```
import dev.langchain4j.mcp.client.McpClient;
```

```

import io.quarkiverse.langchain4j.mcp.runtime.McpClientName;
import jakarta.enterprise.context.ApplicationScoped;
import jakarta.inject.Inject;

@ApplicationScoped
public class McpClientService {

    @Inject
    @McpClientName("helix")
    McpClient helixClient;

    public void useClient() {
        // List tools exposed by the MCP server
        var tools = helixClient.listTools();

        // Execute a tool by name with JSON arguments
        var args = com.fasterxml.jackson.databind.node.JsonNodeFactory.instance
            .objectNode().put("key", "value");
        var request = dev.langchain4j.mcp.client.protocol.McpCallToolRequest.builder()
            .name("some_tool_name")
            .arguments(args)
            .build();
        var result = helixClient.executeTool(request);

        // List and read resources (if the server exposes resources)
        var resources = helixClient.listResources();
        var content = helixClient.readResource("some://resource-uri");
    }
}

```

Use programmatic access when you need to drive tool execution from non-AI code, filter or transform tool results, or integrate MCP resources into your own APIs.

1.3.4. Dev UI: Testing MCP Clients

In development mode (`mvn quarkus:dev`), open the Quarkus Dev UI at <http://localhost:8080/q/dev-ui> and use the **MCP clients** card to:

- View all configured MCP clients
- Browse tools (names, descriptions, parameter schemas)
- Execute tools with custom arguments and inspect results

This is useful for exploring a new MCP server or debugging tool arguments before wiring it into an AI service.

1.3.5. Optional: Resources as Tools

If an MCP server exposes resources, you can expose them as synthetic tools so the LLM can list and read resources:

```
quarkus.langchain4j.mcp.expose-resources-as-tools=true
```

When enabled, the client adds `list_resources` and `get_resource` to the tool set; your AI service can then answer questions like “List available resources” or “Show me the content of resource X” when using `@McpToolBox` for that client.

1.3.6. Optional: Headers and Authorization

- **Static headers:** for HTTP-based clients, set headers via config:

```
quarkus.langchain4j.mcp.helix.header.X-Custom-Header=MyValue
```

- **Dynamic headers:** implement `dev.langchain4j.mcp.client.McpHeadersSupplier` and provide a CDI bean; it will be used for Streamable HTTP clients to add headers per request.
- **Authorization:** implement `io.quarkiverse.langchain4j.mcp.auth.McpClientAuthProvider` to supply a bearer token or API key. Use `@McpClientName` on the bean to restrict it to a specific client. For OIDC, add the dependency `quarkus-langchain4j-oidc-mcp-auth-provider` to propagate the user’s access token to the MCP server.

See the [Quarkus LangChain4j MCP documentation](#) for full details on transports, auth, and the MCP Registry client.

1.4. Agent REST API

The agent layer provides a REST interface at `/api/agent` that mirrors the MCP tools for non-MCP integrations. This is useful for custom agent orchestrators, webhook-based workflows, or testing.

1.4.1. Endpoints

Method	Path	Description
GET	<code>/api/agent/tools</code>	List available gateway tools (optionally filtered by realm)
GET	<code>/api/agent/schema</code>	List all root types (same as <code>query_rootTypes</code>)
GET	<code>/api/agent/schema/{rootType}</code>	JSON Schema-like structure for a single entity type
GET	<code>/api/agent/query-hints</code>	Query grammar summary and example queries
GET	<code>/api/agent/permission-hints</code>	Permission check API summary and examples
POST	<code>/api/agent/execute</code>	Execute a gateway tool by name

1.4.2. Execute Request

The execute endpoint accepts a tool name and arguments:

```
{
  "tool": "query_find",
  "arguments": {
    "rootType": "Location",
    "query": "status:ACTIVE && city:Atlanta",
    "page": { "limit": 10, "skip": 0 }
  }
}
```

```
curl -sS -X POST \
  -H 'Content-Type: application/json' \
  localhost:8080/api/agent/execute \
  -d '{
    "tool": "query_find",
    "arguments": {
      "rootType": "Location",
      "query": "status:ACTIVE && city:Atlanta"
    }
  }'
```

The response shape matches the corresponding Query Gateway REST endpoint (e.g. the Collection envelope for `query_find`).

1.5. Per-Tenant Agent Configuration

Agent behavior can be customized per realm using MicroProfile Config properties:

```
# Run agent tools as a specific user in the "acme" realm
quantum.agent.tenant.acme.runAsUserId=agent-user@acme.com

# Only allow find and plan tools for the "acme" realm
quantum.agent.tenant.acme.enabledTools=query_find,query_plan,query_rootTypes

# Cap find results at 100 for this tenant
quantum.agent.tenant.acme.maxFindLimit=100
```

Configuration properties:

Property	Type	Description
<code>quantum.agent.tenant.<realm>.runAsUserId</code>	string	Optional <code>userId</code> whose security context is used for tool execution
<code>quantum.agent.tenant.<realm>.enabledTools</code>	comma-separated	Optional list of tool names to expose (all six enabled when empty)

Property	Type	Description
<code>quantum.agent.tenant.<realm>.maxFindLimit</code>	integer	Optional maximum number of results for <code>query_find</code>

The default implementation (`PropertyTenantAgentConfigResolver`) reads these from `application.properties` or environment variables. You can replace it by providing a CDI bean implementing `TenantAgentConfigResolver`.

1.5.1. Run-As Principal

When `runAsUserId` is configured, tool execution runs under that user's security context. To enable this, provide a CDI bean implementing `RunAsPrincipalResolver`:

```
import com.e2eq.framework.api.agent.RunAsPrincipalResolver;
import com.e2eq.framework.model.securityrules.PrincipalContext;
import jakarta.enterprise.context.ApplicationScoped;
import java.util.Optional;

@ApplicationScoped
public class MyRunAsPrincipalResolver implements RunAsPrincipalResolver {

    @Override
    public Optional<PrincipalContext> resolvePrincipalContext(String realm, String
userId) {
        // Look up user and build PrincipalContext
        // Return Optional.empty() to fall back to caller's context
    }
}
```

1.6. Agent Configuration API

The framework stores agent configurations as realm-scoped `Agent` entities in MongoDB (`agents` collection). Each agent pairs an LLM reference with a system prompt and a tool filter, enabling different personas for different use cases (e.g., a supply-chain analyst vs. a customer-service bot).

1.6.1. Agent Entity

Field	Type	Description
<code>refName</code>	string	Unique reference name within the realm (e.g., <code>supply-chain-assistant</code>)
<code>name</code>	string	Display name
<code>llmConfigRef</code>	string	Reference to an LLM secret or config (e.g., <code>claude-sonnet</code> , <code>gpt-4o</code>)

Field	Type	Description
<code>context</code>	list of PromptStep	Ordered system/user prompt steps that define the agent's persona
<code>enabledTools</code>	list of string	Tool names this agent can use (e.g., <code>["query_rootTypes", "query_find", "query_plan"]</code>)

Each `PromptStep` has `order` (int), `role` (string: "system" or "user"), and `content` (string: prompt text).

1.6.2. CRUD Endpoints

All endpoints are at `/api/agent/config` and require a `realm` query parameter.

Method	Path	Description
GET	<code>/api/agent/config/list?realm=X</code>	List all agents in the realm
GET	<code>/api/agent/config/{refName}?realm=X</code>	Get an agent by reference name
GET	<code>/api/agent/config/id/{id}?realm=X</code>	Get an agent by ObjectId
POST	<code>/api/agent/config?realm=X</code>	Create or update an agent
DELETE	<code>/api/agent/config/{refName}?realm=X</code>	Delete an agent by reference name
DELETE	<code>/api/agent/config/id/{id}?realm=X</code>	Delete an agent by ObjectId

1.6.3. Example: Creating an Agent


```
curl -sS -X POST \
-H 'Content-Type: application/json' \
'localhost:8080/api/agent/config?realm=acme-corp' \
-d '{
  "refName": "supply-chain-assistant",
  "name": "Supply Chain Assistant",
  "llmConfigRef": "claude-sonnet",
  "context": [
    { "order": 1, "role": "system", "content": "You are a supply-chain analyst. Help
users query orders, track shipments, and monitor inventory." }
  ],
  "enabledTools": ["query_rootTypes", "query_plan", "query_find", "query_count"]
}'
```



Agent configuration complements the property-based tenant configuration ([Per-Tenant Agent Configuration](#)). Properties set operator-level overrides (runAsUserId, maxFindLimit); Agent entities store per-use-case personas and tool filters that can be managed via API.

1.7. Architecture

The following diagram shows how the MCP server, agent REST API, and MCP client relate:



Key design points:

- **Shared execution path:** Both MCP tools and REST agent endpoints delegate to **AgentExecuteHandler**, which routes to **QueryGatewayResource**. Security rules, realm resolution, and query execution are identical regardless of entry point.
- **Zero reverse dependencies:** The **quantum-mcp-server** module depends on **quantum-framework** but the framework has no knowledge of MCP. Applications that do not need MCP simply omit this module.
- **Tenant isolation:** Realm-scoped execution, optional per-tenant tool filtering, and run-as support ensure multi-tenant safety.

1.8. Ontology Tools

In addition to the six gateway tools, the MCP server exposes two ontology discovery tools via

McpOntologyTools:

Tool Name	Description
<code>query_relationships</code>	Find ontology edges (relationships) for an entity — outgoing, incoming, or both, with optional predicate filtering
<code>query_predicates</code>	List ontology predicates (relationship types) defined in the TBox schema, with optional domain/range filtering

These tools allow AI assistants to explore the entity relationship graph. For example, an agent can discover that an `Order` entity has a `placedInOrg` relationship to an `Organization`, then use `query_find` with appropriate filters to retrieve related data.

1.9. Integrating with Claude

Claude supports MCP through Claude Desktop, Claude Code (CLI), and the browser-based Claude.ai (via Custom Connectors). The Quantum MCP server works with all three.

1.9.1. Claude Desktop

Edit `~/Library/Application Support/Claude/claude_desktop_config.json` (macOS) or the equivalent on your platform:

```
{
  "mcpServers": {
    "quantum": {
      "url": "http://localhost:8080/mcp"
    }
  }
}
```

Restart Claude Desktop. The Quantum tools appear in the tool picker and Claude can call them during conversations.

1.9.2. Claude Code (CLI)

Claude Code discovers MCP servers from project-level or user-level configuration.

Project-level — create `.claude/settings.json` in your repository root:

```
{
  "mcpServers": {
    "quantum-mcp": {
      "url": "http://localhost:8080/mcp"
    }
  }
}
```

User-level — edit `~/.claude/settings.local.json` to add the server globally:

```
{
  "mcpServers": {
    "quantum-mcp": {
      "url": "http://localhost:8080/mcp"
    }
  }
}
```

Once configured, Claude Code can use the Quantum tools in any conversation:

> Use `query_rootTypes` to list available entity types

Claude calls `query_rootTypes` and returns:
Location, Order, CodeList, UserProfile, ...

> Find all active locations in the West region

Claude calls `query_find` with:
rootType: "Location"
query: "status:ACTIVE && region:West"
limit: 10

Returns matching Location entities.

1.9.3. Claude.ai (Browser) with Custom Connectors

For cloud-hosted deployments, expose the `/mcp` endpoint over HTTPS and configure a Claude Custom Connector:

1. In Claude.ai, go to Settings > Integrations > Custom Connectors.
2. Add a new connector with your server URL (e.g., <https://api.mycompany.com/mcp>).
3. Configure authentication headers if required by your deployment.

1.9.4. Cursor IDE

Create `.cursor/mcp.json` in your project root:

```
{
  "mcpServers": {
    "quantum": {
      "url": "http://localhost:8080/mcp"
    }
  }
}
```

Or add the server URL in Cursor Settings > MCP Servers.

1.10. How MCP Maps to Framework Agent and Tool Concepts

The framework uses a layered architecture where MCP is one of several access paths to the same underlying tool execution engine. Understanding how these concepts relate helps when extending the framework or building custom integrations.

1.10.1. Concept Mapping

MCP Concept	Framework Concept	Relationship
MCP Tool	@Tool-annotated method	Each MCP tool is a Java method annotated with @Tool (Quarkiverse MCP Server extension) in McpGatewayTools. Tools are code-defined and discovered at startup via classpath scanning — there is no database-backed tool registry.
MCP Tool Name	Method-derived name	The tool name used in MCP tools/call (e.g., query_find) is defined in the @Tool annotation or derived from the method name.
MCP tools/call	AgentExecuteHandler.execute(tool, arguments)	Both the MCP server (McpGatewayTools) and the REST agent endpoint (POST /api/agent/execute) delegate to the same handler.
MCP Resource	SchemaService / QueryHintsProvider	MCP resources (quantum://schema, quantum://query-hints, quantum://permission-hints) are backed by the same provider beans that serve the REST /api/agent/schema and /api/agent/query-hints endpoints.
MCP Client	quarkus-langchain4j-mcp	The framework can also consume tools from external MCP servers. External MCP servers are configured via ToolProviderConfig (connection details) and quarkus.langchain4j.mcp.* properties.
Agent Config	Agent entity	An Agent is a realm-scoped configuration envelope that pairs an LLM reference with a system prompt (PromptSteps) and a tool filter (enabledTools). Managed via /api/agent/config/* CRUD API.

1.10.2. Execution Flow Comparison

Both MCP and REST reach the same execution path:





Security (realm resolution, permission checks, `@FunctionalAction`) is enforced identically at the `QueryGatewayResource` level, regardless of whether the request came through MCP or REST.

1.11. End-to-End Example: Claude Querying a Supply Chain Application

This walkthrough demonstrates a complete flow from starting the application to having Claude interact with live data through MCP.

1.11.1. Prerequisites

- Java 17+, Maven 3.9+
- MongoDB running locally (or via Quarkus Dev Services)
- Claude Desktop or Claude Code installed
- Your Quantum-based application (e.g., a supply chain app with `Location`, `Order`, `Shipment` entities)

1.11.2. Step 1: Add the MCP Server Module

In your application POM, add the `quantum-mcp-server` dependency:

```

<dependency>
  <groupId>com.end2endlogic</groupId>
  <artifactId>quantum-mcp-server</artifactId>
  <version>${quantum.version}</version>
</dependency>
  
```

1.11.3. Step 2: Configure Agent Properties

In `src/main/resources/application.properties`:

```

# --- Authentication ---
auth.provider=custom
quarkus.smallrye-jwt.enabled=true
  
```

```

mp.jwt.verify.publickey.location=publicKey.pem
mp.jwt.verify.issuer=https://myapp.example.com/issuer
mp.jwt.verify.audiences=my-api-client
auth.jwt.secret=${JWT_SECRET:dev-secret-change-in-prod}
auth.jwt.expiration=60
auth.jwt.refresh-expiration=120

# --- Per-tenant agent configuration ---
# The "default" realm: agent runs as a service user, limited to read-only tools
quantum.agent.tenant.default.runAsUserId=agent-svc@mycompany.com
quantum.agent.tenant.default.enabledTools=query_rootTypes,query_plan,query_find,query_count
quantum.agent.tenant.default.maxFindLimit=200

# The "acme-corp" realm: full CRUDL access
quantum.agent.tenant.acme-corp.runAsUserId=bot@acme-corp.com
quantum.agent.tenant.acme-corp.enabledTools=query_rootTypes,query_plan,query_find,query_count,query_save,query_delete

```

1.11.4. Step 3: Start the Application

```
mvn quarkus:dev
```

The MCP endpoint is available at <http://localhost:8080/mcp>.

1.11.5. Step 4: Configure Claude

Claude Desktop — edit `claude_desktop_config.json`:

```

{
  "mcpServers": {
    "supply-chain": {
      "url": "http://localhost:8080/mcp"
    }
  }
}

```

Claude Code — create `.claude/settings.json` in the project root:

```

{
  "mcpServers": {
    "supply-chain": {
      "url": "http://localhost:8080/mcp"
    }
  }
}

```

```
}
```

1.11.6. Step 5: Interact with Claude

Below is a realistic conversation showing how Claude uses the MCP tools:

User: What entity types are available in the supply chain system?

Claude calls `query_rootTypes` and responds:

The system has these entity types: Location, Order, Shipment, Product, Customer, CodeList, UserProfile, Organization.

User: Show me all active warehouses in the Southeast region.

Claude calls `query_find`:

- `rootType: "Location"`
- `query: "status:ACTIVE && type:WAREHOUSE && region:Southeast"`
- `limit: 20`

Found 3 active warehouses in the Southeast:

1. ATL-WH-01 — Atlanta Distribution Center (Atlanta, GA)
2. MIA-WH-03 — Miami Fulfillment Hub (Miami, FL)
3. CLT-WH-02 — Charlotte Regional Warehouse (Charlotte, NC)

User: How many open orders are there for the Atlanta warehouse?

Claude calls `query_count`:

- `rootType: "Order"`
- `query: "status:OPEN && warehouseRef:ATL-WH-01"`

There are 47 open orders for the Atlanta Distribution Center (ATL-WH-01).

User: Show me the last 5 orders with their customer details expanded.

Claude calls `query_plan` first to check expand support:

- `rootType: "Order"`
- `query: "expand(customer) && warehouseRef:ATL-WH-01"`

Then calls `query_find`:

- `rootType: "Order"`

- `query: "expand(customer) && warehouseRef:ATL-WH-01"`
- `limit: 5`

Here are the 5 most recent orders for ATL-WH-01, with customer details:

1. ORD-9921 — Acme Corp (acme@example.com), 12 items, \$2,340
2. ORD-9918 — Beta Industries (beta@example.com), 3 items, \$890 ...

User: What relationships does the Atlanta warehouse have?

Claude calls `query_relationships`:

- `entityId: "ATL-WH-01"`
- `direction: "both"`

The Atlanta warehouse has these relationships:

Outgoing: - `locatedIn` → Organization "Southeast Division" - `servesRegion` → Region "Southeast"

Incoming: - `fulfilledAt` ← 47 open Orders - `shipsFrom` ← 12 active Shipments

1.11.7. What Happened Under the Hood

Claude Desktop / Claude Code

```
|
| (1) MCP tools/list -> discovers 9 tools + 3 resources
|
| (2) MCP resources/read "quantum://query-hints"
|     -> learns BIAPI query syntax, expand(), ontology edges
|
| (3) MCP tools/call "query_rootTypes"
|     -> McpGatewayTools.query_rootTypes()
|         -> AgentExecuteHandler.execute("query_rootTypes", {})
|         -> QueryGatewayResource.listRootTypes()
|
| (4) MCP tools/call "query_find" { rootType, query, limit }
|     -> McpGatewayTools.query_find(rootType, query, realm, limit, skip)
|         -> AgentExecuteHandler.execute("query_find", args)
|         -> QueryGatewayResource.find(FindRequest)
|             -> Morphia Query -> MongoDB
|
| (5) MCP tools/call "query_relationships" { entityId, direction }
|     -> McpOntologyTools.query_relationships(entityId, direction, ...)
|         -> OntologyEdgeRepo.findBySrc() / findByDst()
|             -> MongoDB edges collection
|
```

V
JSON response serialized and returned to Claude

At every step, the Quantum security filter validates the caller's identity, resolves the realm, and applies permission rules. The agent never bypasses `@FunctionalAction` checks or data scoping.

1.11.8. Adding Ontology Context for Richer Queries

When your application has an ontology TBox defined (see [Ontologies in Quantum](#)), Claude can use the `query_predicates` tool to discover relationship types and then use `query_relationships` to traverse the graph.

For example:

User: What types of relationships exist between Orders and Organizations?

Claude calls `query_predicates`:

- `domainFilter: "Order"`
- `rangeFilter: "Organization"`

Found 2 predicates:

1. `placedByOrg` — Order → Organization (the org that placed the order)
2. `fulfilledByOrg` — Order → Organization (the org fulfilling the order)

This lets Claude build more targeted queries and provide richer answers about entity relationships.

1.12. MCP Server Configuration Reference

The MCP server requires no additional configuration beyond including the `quantum-mcp-server` module. The `/mcp` endpoint is automatically registered by the Quarkus MCP Server extension.

Optional properties:

Property	Default	Description
<code>quarkus.mcp.server.traffic-logging</code>	<code>NONE</code>	Set to <code>ALL</code> to log MCP JSON-RPC messages for debugging
<code>feature.queryGateway.execution.enabled</code>	<code>false</code>	Set to <code>true</code> to enable <code>expand()</code> (AGGREGATION mode) in <code>query_find</code>
<code>quantum.queryGateway.deleteMany.maxMatches</code>	<code>2000</code>	Maximum entities that <code>query_deleteMany</code> will delete in a single call

See [AI Agent Integration](#) for the full agent API design and additional scenarios.